

Online Result Summary


Model: selfie_segmentation

Constraint targets: Max p99 Latency: 50 ms


In 16 measurements across 4 configurations, **selfie_segmentation_config_2** is **266%** better than the default configuration at meeting the objectives, under the given constraints.

- **selfie_segmentation_config_2**: 4 CPU instances with a max batch size of 0 on platform python

Curves corresponding to the 3 best model configuration(s) out of a total of 4 are shown in the plots.



Throughput vs. Latency curves for 3 best configurations.



CPU Memory vs. Latency curves for 3 best configurations.

The following table summarizes each configuration at the measurement that optimizes the desired metrics under the given constraints.

Model Config Name	Max Batch Size	Dynamic Batching	Total Instance Count	p99 Latency (ms)	Throughput (infer/sec)	Max CPU Memory Usage (MB)
selfie_segmentation_config_2	0	Disabled	4:CPU	20.074	494.004	53.0
selfie_segmentation_config_1	0	Disabled	2:CPU	16.606	265.468	48.0
selfie_segmentation_config_0	0	Disabled	1:CPU	32.426	138.086	47.0
selfie_segmentation_config_default	0	Disabled	1:CPU	17.455	134.952	46.0